

Performance Infrastructure Improvements in z/OS V2.5 WLM

Tuesday
August 10th, 2021

Andreas Henicke, z/OS Workload Management
andreas.henicke@de.ibm.com
IBM Germany Research & Development, 2021

Trademarks



Not all common law marks used by IBM are listed on this page. Failure of a mark to appear does not mean that IBM does not use the mark nor does it mean that the product is not actively marketed or is not significant within its relevant market.

Those trademarks followed by ® are registered trademarks of IBM in the United States; all others are trademarks or common law marks of IBM in the United States.

For a complete list of IBM Trademarks, see www.ibm.com/legal/copytrade.shtml:

*, AS/400®, e business(logo)®, DBE, ESCO, eServer, FICON, IBM®, IBM (logo)®, iSeries®, MVS, OS/390®, pSeries®, RS/6000®, S/30, VM/ESA®, VSE/ESA, WebSphere®, xSeries®, z/OS®, zSeries®, z/VM®, System i, System i5, System p, System p5, System x, System z, System z9®, BladeCenter®

The following are trademarks or registered trademarks of other companies.

Adobe, the Adobe logo, PostScript, and the PostScript logo are either registered trademarks or trademarks of Adobe Systems Incorporated in the United States, and/or other countries.

Cell Broadband Engine is a trademark of Sony Computer Entertainment, Inc. in the United States, other countries, or both and is used under license therefrom.

Java and all Java-based trademarks are trademarks of Sun Microsystems, Inc. in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

Intel, Intel logo, Intel Inside, Intel Inside logo, Intel Centrino, Intel Centrino logo, Celeron, Intel Xeon, Intel SpeedStep, Itanium, and Pentium are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

ITIL is a registered trademark, and a registered community trademark of the Office of Government Commerce, and is registered in the U.S. Patent and Trademark Office.

IT Infrastructure Library is a registered trademark of the Central Computer and Telecommunications Agency, which is now part of the Office of Government Commerce.

* All other products may be trademarks or registered trademarks of their respective companies.

Notes:

Performance is in Internal Throughput Rate (ITR) ratio based on measurements and projections using standard IBM benchmarks in a controlled environment. The actual throughput that any user will experience will vary depending upon considerations such as the amount of multiprogramming in the user's job stream, the I/O configuration, the storage configuration, and the workload processed. Therefore, no assurance can be given that an individual user will achieve throughput improvements equivalent to the performance ratios stated here.

IBM hardware products are manufactured from new parts, or new and serviceable used parts. Regardless, our warranty terms apply.

All customer examples cited or described in this presentation are presented as illustrations of the manner in which some customers have used IBM products and the results they may have achieved. Actual environmental costs and performance characteristics will vary depending on individual customer configurations and conditions.

This publication was produced in the United States. IBM may not offer the products, services or features discussed in this document in other countries, and the information may be subject to change without notice. Consult your local IBM business contact for information on the product or services available in your area.

All statements regarding IBM's future direction and intent are subject to change or withdrawal without notice, and represent goals and objectives only.

Information about non-IBM products is obtained from the manufacturers of those products or their published announcements. IBM has not tested those products and cannot confirm the performance, compatibility, or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products.

Prices subject to change without notice. Contact your IBM representative or Business Partner for the most current pricing in your geography.

Notice Regarding Specialty Engines (e.g., zIIPs, zAAPs and IFLs):

Any information contained in this document regarding Specialty Engines ("SEs") and SE eligible workloads provides only general descriptions of the types and portions of workloads that are eligible for execution on Specialty Engines (e.g., zIIPs, zAAPs, and IFLs). IBM authorizes customers to use IBM SE only to execute the processing of Eligible Workloads of specific Programs expressly authorized by IBM as specified in the "Authorized Use Table for IBM Machines" provided at www.ibm.com/systems/support/machine_warranties/machine_code/aut.html ("AUT").

No other workload processing is authorized for execution on an SE.

IBM offers SEs at a lower price than General Processors/Central Processors because customers are authorized to use SEs only to process certain types and/or amounts of workloads as specified by IBM in the AUT.

Agenda

- Batch Initiator Enhancements for Job Placement and Initiator Management
- Service Coefficient Removal
- System Recovery Boost
- Other Enhancements and Important Information

Batch Initiator Enhancements for Job Placement and Initiator Management

WLM Batch Initiator Management

Policy adjustment:


- Initiators are dynamically started by WLM to meet service class period goals by reduction of batch queue delays
- WLM selects the system considering available system resources based on the job's importance

Resource adjustment:


- Start additional initiators if jobs are waiting and unused capacity is available
 - Based on service class history
 - Up to 5 at a time

Problem:

Initiator placement is determined based on the available or displaceable capacity of **general purpose processors**. As the capacity of specialty processors is increasing, as does the demand for such workloads, it is important to consider the workload specific capacity demand for placement routines



IBM z/OS V2R5 Batch Initiator Enhancements



Starting with IBM z/OS V2R5, WLM Batch Initiator takes the service class consumption of specialty processors into account.

Initiators for batch jobs classified into a service class of which the first period executes primarily on zIIP will tend to be added on systems with spare zIIP capacity.

Initiators for batch jobs classified into a service class of which the first period executes primarily on CP will tend to be added on systems with spare CP capacity.

This enhancement will help to remove the need of manual management of zIIP eligible batch jobs such that the z/OS system automatically selects to best suited system.

Mixed Service Classes

Mixed service classes:

Initiators running CP workload and initiators running zIIP workload are classified into the same service class. WLM then must assume that on such system each initiator consumes processor capacity on both processor types. Initiators will only be started if there is sufficient available capacity for all checked processor types.

To allow the customer some control with “mixed” service classes, WLM does the following:

- The available CP capacity is always checked.
- If OPT parameter “IIPHONORPRIORITY = NO” is in effect, the available zIIP capacity is always checked.
- If OPT parameter “IIPHONORPRIORITY = YES” is in effect (the default), zIIP capacity is only checked if the service class’ consumed amount of zIIP capacity is at least ten percent of the consumed service units on CPs

Batch Initiator Enhancements Recommendation



- Use separate sets of service classes for initiators running on specialty processors and initiators executing on standard CPs such that all initiators running in the same service class either execute on standard CPs or on specialty processors. This is no mandatory requirement, but it allows for a more accurate management of the initiators.
- If back-level systems and IBM z/OS V2R5 systems in a SYSPLEX/MAS have WLM managed initiators running in the same service class and on either system the initiators execute on specialty processors, balancing these initiators across the SYSPLEX/MAS will not be optimal. Neither system would be aware of the other system's consumption of specialty processor capacity.
- You can use one set of service classes for initiators executing zIIP workloads on IBM z/OS V2R5 systems and another set of service classes for initiators executing zIIP workloads on back-level systems.

Service Coefficient Removal

Removal of Service Coefficients

```

Coefficients/Options  Notes  Options  Help
-----
      Service Coefficient/Service Definition Options
Command ===> -----

Enter or change the Service Coefficients:

CPU   . . . . . 1.0      (0.1-99.9)
IOC   . . . . . 0.0      (0.0-99.9)
MSO   . . . . . 0.0000   (0.0000-99.9999)
SRB   . . . . . 1.0      (0.0-99.9)

Enter or change the service definition options:

I/O priority management . . . . . YES (Yes or No)
Enable I/O priority groups . . . . . YES (Yes or No)
Dynamic alias tuning management . . . . . YES (Yes or No)
Deactivate Discretionary Goal Management YES (Yes or No)

```

Statement of Direction

z/OS V2.4 is the last release of z/OS to allow specifying service coefficients in the Workload Manager (WLM) service definition on the Service Definition Details page.

Starting with z/OS V2.5 the IBM recommended values CPU=1, SRB=1, MSO=0, and IOC=0, will be the predefined values.

IBM recommends that you adjust your service coefficients before upgrading to a later release.

New health checks ZOSMIGV2R4_NEXT_WLM_ServCoeff and ZOSMIGV2R3_NEXT2_WLM_ServCoeff are provided to check the WLM service coefficients being specified in the currently installed WLM service definition.

Migration Health Check

- OA59066 provides two health checks
 - ZOSMIGV2R3_NEXT2_WLM_ServCoeff at z/OS 2.3
 - ZOSMIGV2R4_NEXT_WLM_ServCoeff at z/OS 2.4
 - As is customary with migration health check, they are inactive and must be activated by the customer explicitly
- If the health check is activated and the recommended WLM service coefficients are already in use, information message **IWMH101I THE RECOMMENDED WLM SERVICE COEFFICIENTS ARE BEING USED** is issued.
- If the health check is activated and one or more of the WLM service coefficients currently in use have a different value than CPU=1, SRB=1, IOC=0, MSO=0, exception message **IWMH102E THE RECOMMENDED WLM SERVICE COEFFICIENTS ARE NOT BEING USED** is issued.

Overview – Migration Health Check

- If exception message IWMH102E is issued, the customer should consider changing the service coefficients in the WLM service definition to the recommended values and review the durations for any multi-period service classes and adjust them accordingly.
- Information how to adjust the durations is provided in the APAR text
- z/OS MVS Planning: Workload Management Chapter 13 (Defining service coefficients and options) is updated accordingly

Duration (DUR) is currently calculated as:

$$\text{DUR} = (\text{CPU} * \text{CPU service units}) + (\text{SRB} * \text{SRB service units}) + (\text{I/O} * \text{I/O service units}) + (\text{MSO} * \text{storage service units})$$

where CPU, SRB, I/O, and MSO are the installation defined WLM service coefficients. With CPU=1, SRB=1, I/O=0, MSO=0 the new duration is simply calculated as:

$$\text{NEW DUR} = \text{CPU service units} + \text{SRB service units}$$

Converting DUR into NEW DUR is calculated as:

$$\text{NEW DUR} = \text{DUR} / \text{Total service units} * (\text{CPU service units} / \text{CPU} + \text{SRB service units} / \text{SRB})$$

where CPU and SRB are the old service coefficients and Total service units is the sum of CPU, SRB, I/O, and MSO service units. These values should be collected for a peak period interval from, for example, the RMF Postprocessor Workload Activity (WLMGL) report.

EXAMPLE:

- Old DUR is 90000
- Old default service coefficients used (CPU=10, SRB=10)
- Values from RMF WLMGL peak period interval:
TOTAL_SU = 6218K
CPU_SU = 5877K
SRB_SU = 95667

$$\text{NEW DUR} = 90000 / 6218K * (5877K / 10 + 95667 / 10) = 8645$$

SDSF Health Checker Display (SDSF CK)

The initial state of the migration health check is **INACTIVE**.

```
SDSF HEALTH CHECKER DISPLAY  SYS1                LINE 1-1 (1)
COMMAND INPUT ==>                                SCROLL ==>
PREFIX=IS*  DEST=(ALL)  OWNER=*  SORT=ModName/A  SYSNAME=SYS1  FILTERS=1
NP  NAME                                CheckOwner  State
    ZOSMIGV2R4_NEXT_WLM_SERVCOEFF      IBMWLM      INACTIVE(ENABLED)
```

Check is activated by line command **A** in CK display (**H** would deactivate the check).

```
SDSF HEALTH CHECKER DISPLAY  SYS1                LINE 1-1 (1)
COMMAND INPUT ==>                                SCROLL ==>
PREFIX=IS*  DEST=(ALL)  OWNER=*  SORT=ModName/A  SYSNAME=SYS1  FILTERS=1
NP  NAME                                CheckOwner  State
    ZOSMIGV2R4_NEXT_WLM_SERVCOEFF      IBMWLM      ACTIVE(ENABLED)
```

```
CHECK(IBMWLM,ZOSMIGV2R4_NEXT_WLM_SERVCOEFF)
SYSPLEX:   PLEX1      SYSTEM: SYS1
START TIME: 12/15/2020 06:12:47.798322
CHECK DATE: 20200101  CHECK SEVERITY: LOW

* Low Severity Exception *

IWMH102E The recommended WLM service coefficients are not being used.

Explanation:  z/OS V2R4 is the last release to allow specifying
service coefficients in the WLM service definition. With the release
after z/OS V2R4, the coefficients are pre-set to CPU=1, IOC=0,
MS0=0, and SRB=1; and cannot be changed.

One or more of the WLM service coefficients currently in use have a
different value and will be ignored after z/OS V2R4:

CPU=10.0 IOC=5.00 MS0=3.0000 SRB=10.0

These values are different because you have specified them in your
currently installed WLM service definition or defaults are taken if
you did not define them.

System Action:  Processing continues.

Operator Response:  Report this problem to the system programmer.

System Programmer Response:  Change the service coefficients in your
WLM service definition (CPU=1, IOC=0, MS0=0, SRB=1) and review the
durations for any multi-period service classes and adjust them
accordingly. Also, you might need to adjust your accounting
procedures.
```

Line command **SB** displays the result of the check.

The WLM service coefficients currently in use are part of the message explanation.

z/OS Console

Check can also be activated by MODIFY console command:
f hzsproc,activate,check=(ibmwlm,*)

```
- SYS1 f hzsproc,activate,check=(ibmwlm,*)
SYS1 HZS0400I CHECK(IBMWLM,*):
ACTIVATE PROCESSING HAS BEEN COMPLETED
| SYS1 HZS0001I CHECK(IBMWLM,ZOSMIGV2R4_NEXT_WLM_SERVCOEFF):
| IWMH102E The recommended WLM service coefficients are not being used.
```

Important note:

Customers typically don't need to care about activating migration health checks. When doing an upgrade, the **z/OSMF Upgrade Workflow** will run the migration health checks for the customers automatically.

Coexistence with z/OS V2.5

- z/OS V2.5 uses pre-set WLM service coefficients. They can no longer be defined in the WLM ISPF Application and the z/OSMF WLM task.
- WLM must take care that all systems in the sysplex use identical service coefficients.
 - When the WLM service definition installed has no service coefficients defined, defaults (CPU=10, SRB=10, IOC=5, MSO=0) are used, or when the service definition specifies other coefficients than the recommended ones, the V2.5 systems would use other coefficients than the pre-V2.5 systems.
 - Of course, this should not happen when the customer properly prepared the V2.5 migration. But it could happen, so we must deal with it...
 - Thus, coexistence logic is available which checks for a mismatch of the service coefficients and internally applies identical values.

Adapt Your Multiperiod Durations

- If the customer did not prepare his WLM service definition for the removal of the service coefficients, following steps should be taken because the calculation of DURATION for multi-period service classes changes:

Before z/OS V2.5 the DURATION is calculated as:

$OLD\ DUR = (CPU * CPU\ service\ units) + (SRB * SRB\ service\ units) + (IOC * I/O\ service\ units) + (MSO * storage\ service\ units)$

where CPU, SRB, IOC, and MSO are the installation defined WLM service coefficients. With CPU=1, SRB=1, IOC=0, MSO=0 the new duration is simply calculated as:

$NEW\ DUR = CPU\ service\ units + SRB\ service\ units$

Converting OLD DUR into NEW DUR is calculated as:

$NEW\ DUR = OLD\ DUR / Total\ service\ units * (CPU\ service\ units / CPU + SRB\ service\ units / SRB)$

where CPU and SRB are the old service coefficients and Total service units is the sum of CPU, SRB, IOC, and MSO service units. CPU, SRB, and Total service unit values should be collected for a peak period interval from, for example, the RMF Postprocessor Workload Activity (WLMGL) report.

EXAMPLE: $OLD\ DUR = 90000$

- Old default service coefficients used (CPU=10, SRB=10)
- Values from RMF WLMGL peak period interval:
 - TOTAL_SU = 6218K
 - CPU_SU = 5877K
 - SRB_SU = 95667

$NEW\ DUR = 90000 / 6218K * (5877K / 10 + 95667 / 10) = 8645$

Copyright© by SHARE Association Except where otherwise noted, this work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivs 3.0 license. <http://creativecommons.org/licenses/by-nc-nd/3.0/>

System Recovery Boost


Introduction

- Part of the IBM Z Instant Recovery offering
- Minimizes impact and duration of downtimes
- No additional MSU consumption or IBM software licensing costs

“Because planned downtime is costly, some enterprises try to avoid it for needed bug fixes and security updates... [this] results in a greater risk of unplanned downtime, creating a downtime cycle.”

Forrester Consulting, 2019

Anatomy of a planned outage



Technical Overview

Speed Boost

- Subcapacity CPC Boost
- Upgrade to full capacity on activate
- Affects **only** active LPAR

zIIP Boost

- Enables zIIPs to process non-system CP workload
- Activates additional zIIPs to the LPAR
- **May** affect LPARs on same CPC

Technical Overview

- Intelligent Resource Director (IRD) weight management is deactivated
- (Tenant) Resource Group Capping and Discretionary Goal Management (DGM) are deactivated
- Routing changes
 - Routing recommendations change according to boosted capacity (**IPL only**)
 - Deactivated (**Shutdown only**)

How to benefit from System Recovery Boost

- IBM z15
- IEASYSxx **BOOST** parameter
BOOST=SYSTEM | ZIIP | SPEED | NONE
- **HIPERDISPATCH=YES** (zIIP Boost only)
with at least 1 shared zIIP



IPL Boost Processing

```
SYS1  SHASP100 IEABE      ON STCINRDR  
SYS1  SHASP373 IEABE      STARTED  
SYS1  IEA678I  All IPL boosts have ended  
SYS1  SHASP395 IEABE      ENDED - RC=0000  
SYS1  IWM063I WLM POLICY WAS REFRESHED.  
SYS1  IWM064I BOOST ENDED.
```

1. Automatically started during IPL
 - Signaled via console message and ENF signal (84)
 - Ends current and starts new SMF interval

```
SYS1  IEA675I IPL zIIP boost is active with 0 transient zIIP cores  
SYS1  IRA860I HIPERDISPATCH MODE IS NOW ACTIVE
```

2. Boost Processing Phase
 - Possible offline zIIPs are configured online
 - Non-system work is balanced between zIIPs and CPs (**zIIP Boost**)
 - CPs are configured full capacity (**Speed Boost**)
 - IRD weight management, (T)RG capping and DGM are deactivated
 - Routing weights changes until 3 min before end (configurable via OPT **RTBELeadTime**)
3. Boost End (after **60 min** or **PROC IEABE**)
 - Signaled via console message and ENF signal (84)
 - Ends current and starts new SMF interval
 - Return to normal system behavior

Shutdown Boost Processing

1. Started via **PROC IEASDBS**

- ...

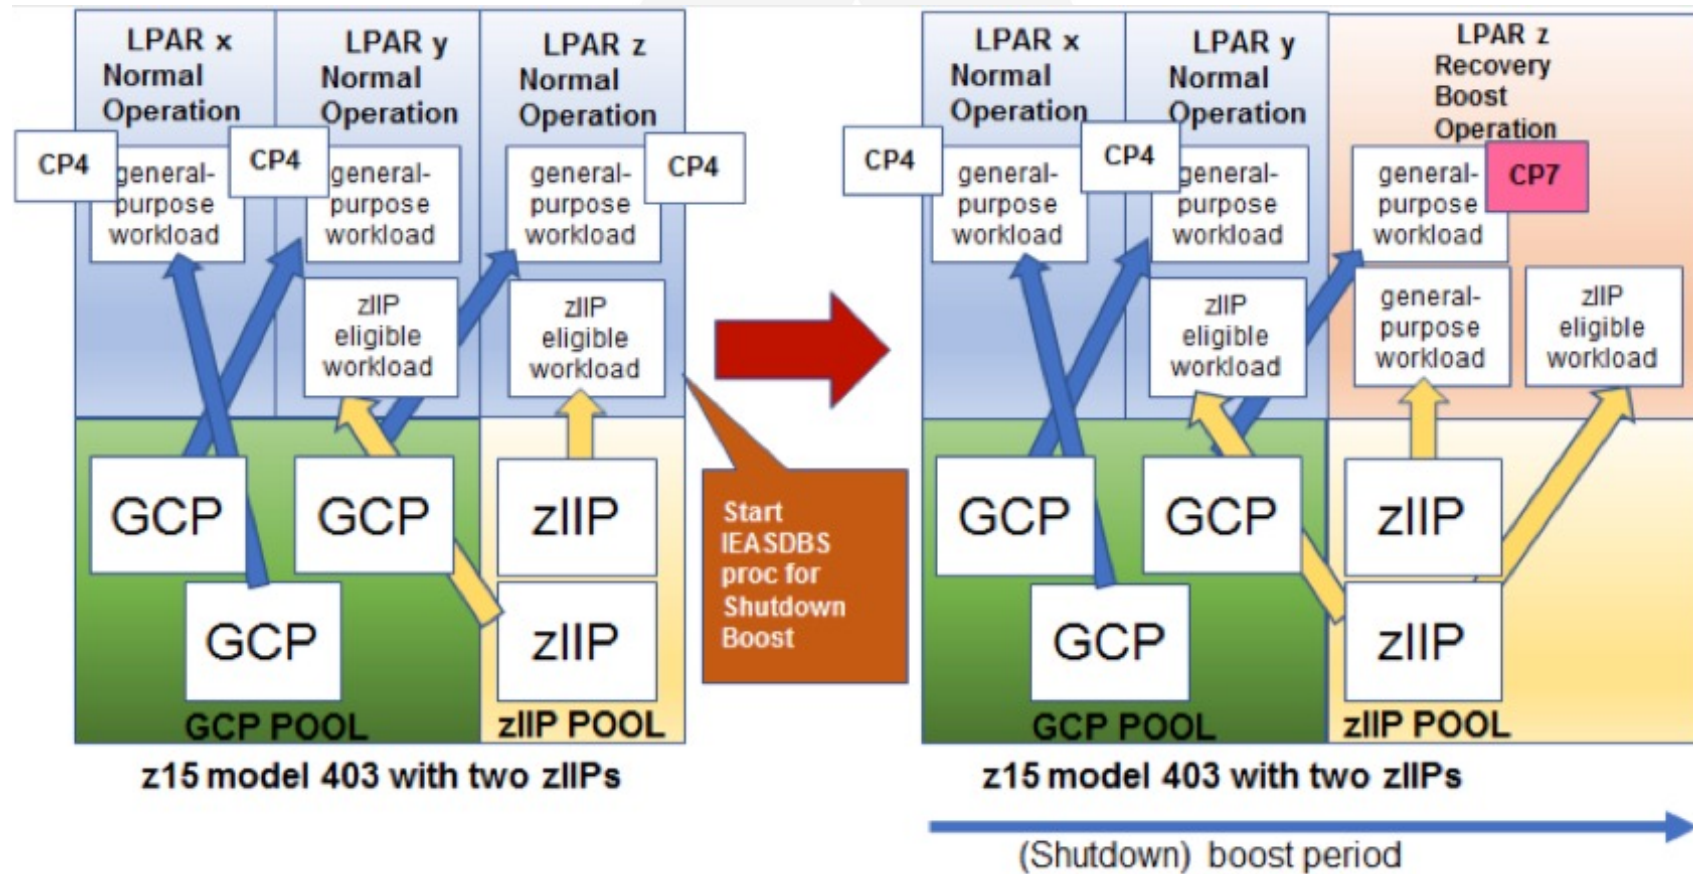
2. Boost Processing Phase

```
SYS1  $HASP100 IEASDBS  ON STCINRDR
SYS1  $HASP373 IEASDBS  STARTED
SYS1  IEA681I Shutdown speed boost is active
SYS1  IEA675I Shutdown zIIP boost is active with 0 transient zIIP cores
SYS1  $HASP395 IEASDBS  ENDED - RC=0000
SYS1  IWM063I WLM POLICY WAS REFRESHED.
SYS1  IWM064I BOOST ACTIVATED.
```

3. Boost End (after 30 min or **PROC IEABE**)

```
SYS1  $HASP100 IEABE    ON STCINRDR
SYS1  $HASP373 IEABE    STARTED
SYS1  IEA678I All shutdown boosts have ended
SYS1  $HASP395 IEABE    ENDED - RC=0000
SYS1  IWM063I WLM POLICY WAS REFRESHED.
SYS1  IWM064I BOOST ENDED.
```

CPC illustration



Capping

Capping Type	CP Type	Interactions with Speed Boost	Interactions with zIIP Boost
LPAR Absolute	CP	Cap Respected	N/A
LPAR Absolute	zIIP	N/A	Cap Ignored
Group Absolute	CP	Cap Respected	N/A
Group Absolute	zIIP	N/A	Cap Ignored
Initial	CP	Cap Respected	N/A
Initial	zIIP	N/A	Cap Ignored
Defined Capacity	CP	Cap Respected	Work running on zIIPs due to blurring doesn't count toward CP Cap; CP Cap respected
Defined Capacity	zIIP	N/A	N/A
LPAR Group Capacity	CP	Cap Respected	Work running on zIIPs due to blurring doesn't count toward CP Cap; CP Cap respected
LPAR Group Capacity	zIIP	N/A	N/A
Resource Group/Tenant Resource Group	CP	Caps Ignored	Caps Ignored
Resource Group/Tenant Resource Group	zIIP	N/A	Caps Ignored

Performance Considerations

- Two main phases during system startup
 - Single processor phase
 - Switch to multiprocessor processing
- Speed Boost does not affect other LPARs
- zIIP Boost can affect other LPARs
 - Considerations regarding weight adjustment needed

Considerations for Automation

- Any automation product, IBM System Automation or vendor provided, can be used
- Possible actions for System Recovery Boost
 - Call PROC IEASDBS
 - Increasing/Decreasing the level of parallelism
 - Activation/Deactivation of temporary capacity record
 - Dynamic LPAR weight modification
 - General automation on issued z/OS messages

Messages

- **IRA867I HIPERDISPATCH=NO CURRENTLY NOT ALLOWED**
- **IWM063I changed to WLM POLICY WAS REFRESHED**
- **IWM064I additions**
 - **IWM064I Boost Activated**
 - **IWM064I Boost Ended**

Monitoring

WLM

- IWMWSYSQ WLM service
SYSI_BoostInfo
- QVS Sysevent
QVSBoostInfo

RMF

- Postprocessor CPU Activity and Workload Activity reports (current LPAR)
- Monitor II CPC and Postprocessor Partition Data reports (LPARs in CPC)

RMF Postprocessor CPU Activity Report

CPU ACTIVITY									
z/OS V2R4		SYSTEM ID R75		DATE 08/07/2019		INTERVAL 12.43.			
CPU 8561		CPC CAPACITY 3068		RPT VERSION V2R4 RMF		TIME 02.02.16		CYCLE 1.000 SEC	
MODEL 634				SEQUENCE CODE 000000000002B0E8					
H/W MODEL T01		CHANGE REASON=NONE		HIPERDISPATCH=YES					
				BOOST TYPE=ALL		BOOST CLASS=IPL			
---CPU---									
NUM	TYPE	ONLINE	LPAR BUSY	MVS BUSY	PARKED	PROD	UTIL	LOG PROC	--I/O INTERRUPTS--
0	CP	100.00	14.89	14.82	0.00	100.00	14		
1	CP	100.00	6.83	6.80	0.00	100.00	6		
2	CP	100.00	8.90	8.84	0.00	100.00	8		
3	CP	100.00	5.62	5.60	0.00	100.00	5		
4	CP	100.00	4.54	4.54	0.00	100.00	4		
5	CP	100.00	2.47	2.47	0.00	100.00	2		
TOTAL/AVERAGE			7.21	7.18		100.00	7		
6	IIP	100.00	7.49	7.48	0.00	100.00	7		
7	IIP	100.00	3.30	3.30	0.00	100.00	3		
8	IIP	100.00	0.00	-----	100.00	100.00	0		
9	IIP	100.00	0.00	-----	100.00	100.00	0		
A	IIP	100.00	0.00	-----	100.00	100.00	0		
B	IIP	100.00	0.00	-----	100.00	100.00	0		
C	IIP	100.00	0.00	-----	100.00	100.00	0		
D	IIP	100.00	0.00	-----	100.00	100.00	0.00	0.0	LOW
E	IIP	100.00	0.00	-----	100.00	100.00	0.00	0.0	LOW
F	IIP	100.00	0.00	-----	100.00	100.00	0.00	0.0	LOW
TOTAL/AVERAGE			1.08	5.39		100.00	1.08	166.7	
----- MULTI-THREADING ANALYSIS -----									
CPU TYPE	MODE	MAX CF	CF	AVG TD					
CP	1	1.000	1.000	1.000					
IIP	1	1.000	1.000	1.000					

BOOST TYPE

The boost type that was active at the end of the interval

NONE

Boost is inactive

ZIIP

zIIP capacity boost

SPEED

Speed boost

ALL

zIIP capacity and speed boost are b

BOOST CLASS:

Indicates whether boost was active during Shutdown. NONE is displayed if boost was inactive.

BOOST TYPE The boost type that was active at the end of the interval

NONE Boost is inactive

ZIIP zIIP capacity boost

SPEED Speed boost

ALL zIIP capacity and speed boost are both active

BOOST CLASS: Indicates whether boost was active during IPL or Shutdown. NONE is displayed if boost was inactive.

RMF Postprocessor Partition Data Report

P A R T I T I O N D A T A R E P O R T												PAGE 3									
z/OS V2R4			SYSTEM ID R75			DATE 08/07/2019			INTERVAL 12.43.749												
			RPT VERSION V2R4 RMF			TIME 02.02.16			CYCLE 1.000 SECONDS												
MVS PARTITION NAME			R75		PHYS PROC NUM		54		GROUP NAME		N/A		INITIAL CAP		NO						
IMAGE CAPACITY			541		CP		34		LIMIT		N/A		LPAR HW CAP		NO						
NUMBER OF CONFIGURED PARTITIONS			44		ICF		10		AVAILABLE		N/A		HW GROUP CAP		YES						
WAIT COMPLETION			NO		IIP		10						ABS MSU CAP		YES						
DISPATCH INTERVAL			DYNAMIC																		
----- PARTITION DATA -----														-- LOGICAL PARTITION PROCESSOR DATA --				-- AVERAGE PROCESSOR UTILIZATION PERCENTAGES --			
-----MSU-----														--CAPPING---				---			
NAME	S	BT	WGT	DEF	ACT	DEF	W	BT							EFFECTIVE		TOTAL				
R75	A	S	500	0	38	NN	Y	S							1.24		1.27				
S50	A	N	500	0	152	NN	N								4.94		4.99				
S51	A	N	500	0	144	NN	N								4.71		4.77				
R76	A*	S	459*	0	32	N*N	Y*	I							1.06		1.13				
S58	A	N	500	0	150	NN	N								4.90		4.95				
S59	A	N	500	0	151	NN	N								4.94		4.99				
PHYSICAL																		0.25			
TOTAL																		21.79		22.36	
XACFT87	A		DED											10.00		10.00					
X5CFT87	A		DED											60.00		60.00					
X7CFT87	A		DED											30.00		30.00					
PHYSICAL																		0.00			
TOTAL																		100.0		100.0	
R75	A	I	500			NN	Y		10	IIP	00.01.21.881	00.01.22.452	1.07	1.08	0.01	1.07	1.08				
S50	A	N	500			NN	N		5	IIP	00.00.28.080	00.00.28.476	0.74	0.75	0.01	0.37	0.37				
S51	A	N	500			NN	N		5	IIP	00.00.27.696	00.00.28.111	0.73	0.74	0.01	0.36	0.37				
R76	A*	I	459*			N*	Y*Y*		10*	IIP	00.00.03.529	00.00.03.583	0.08	0.08	0.00	0.05	0.05				
S58	A	N	500			NN	N		4	IIP	00.00.27.968	00.00.28.239	0.92	0.92	0.00	0.37	0.37				
S59	A	N	500			NN	N		4	IIP	00.00.27.989	00.00.28.396	0.92	0.93	0.01	0.37	0.37				
PHYSICAL																		0.03		0.03	

BT The boost type that was active at some point within the interval

S Speed boost. Only shown for partitions that are grouped together for processor type CP.

I zIIP capacity boost. Only shown for partitions that are grouped together for processor type IIP.

N Speed boost or zIIP capacity boost was inactive

- zIIP capacity boost was inactive if shown for processor type IIP
- Speed boost was inactive if shown for processor type CP

RMF Postprocessor Workload Activity Report

W O R K L O A D A C T I V I T Y

z/OS V2R4

SYSplex SVPlex7
RPT VERSION V2R4 RMF

DATE 08/07/2019
TIME 03.30.00

INTERVAL 14.59.980

POLICY ACTIVATION DATE/TIME 05/14/2019 10.14.45

- SERVICE POLICY PAGE -

SERVICE DEFINITION: SYSTEST BASEPOL from ...DEF.IOSPLEX
INSTALL DATE: 05/14/2019 10.14.44 INSTALLED BY: SETUP
POLICY: BASEPOL BASEPOL from ...DEF.IOSPLEX
DISCRETIONARY GOAL MANAGEMENT: YES
DYNAMIC ALIAS MANAGEMENT: YES
I/O PRIORITY MANAGEMENT: YES

SYSTEMS

---ID---	OPT	SU/SEC	CAP%	--TIME--	INTERVAL	--BOOST--
R75	AA	53156.1	171	03.30.00	00.14.59	A Shutdown

BOOST

Indicates whether boost was inactive or active during IPL or Shutdown.

If active, the boost type is shown

I
S
A

zIIP capacity boost

Speed boost

zIIP capacity and speed boost were both active

RMF Monitor III CPC Report Postprocessor Workload Activity Report

RMF V2R4 CPC Capacity Line 1 of 23
Command ==> _ Scroll ==> CSR

Samples: 60 System: R75 Date: 08/07/19 Time: 02.04.00 Range: 60 Sec

Partition: R75 8561 Model 634 Boost: All
CPC Capacity: 3068 Weight % of Max: 55.6 4h Avg: 1
IIP: 100 AbsMSUCap: Y Group: N/A Limit: N/A

Proc Num	Logical Effect	Util % Total	- Physical LPAR	Util % Effect	- Util % Total
132			0.5	20.5	21.0
6.0	5.7	6.0	0.0	1.0	1.1
6.0	0.0	0.0	0.0	0.0	0.0
20.0	8.3	8.4	0.0	4.9	4.9
20.0	8.3	8.4	0.1	4.9	4.9
50.0	3.3	3.4	0.1	4.9	5.0
30.0	5.5	5.5	0.1	4.8	4.9
PHYSICAL			0.3		0.3
*ICF		10.0	0.0	100	100
XACFT87	N N N	1.0	100	100	10.0
X5CFT87	N N N	6.0	100	100	60.0

BOOST The boost type that was active at the end of MINTIME

N Boost was inactive

zIIP zIIP capacity boost

Speed Speed boost

All zIIP capacity and speed boost were both active

+ hidden fields

Enablement of System Recovery Boost for IBM Z

OA57849 makes sure that all required SRB APARs are installed

	z/OS V2.4	z/OS V2.3
z/OS	OA57849	OA57849
RMF	OA56683	OA56683
CPM	OA57552	OA57552
CIM	OA57478	OA57478

Wrap-Up

- Activated by default with z15
- Two Boost periods per IPL (IPL 60 min, Shutdown 30 min)
- Two Boost types (Speed and zIIP)
- Shutdown Boost has permanent changes for WLM Routing
- No additional costs (4HRA is not increased)
- Understanding and adjusting IPL and Shutdown process can dramatically increase Boost benefit

IBM Hardware Announcement 120-050 – August 4th 2020

IBM z15 Model T01 and IBM z15 Model T02 are enhanced to provide increased security, resiliency, performance, and flexibility for mission-critical workloads in a hybrid cloud

- ...
- Enhanced System Recovery Boost capabilities enable clients to leverage a new class of boost that can be applied to a range of sysplex recovery processes, including sysplex partitioning, CF structure recovery, CF data sharing member recovery, and IBM HyperSwap® recovery
- ...

System Recovery Boost Stage 2

Recovery Process Boost is a boost class that can help accelerate system recovery.

Recovery Process boost periods are restricted to 5 minute durations and are limited to 30 total minutes per partition per 24 hour period.

Recovery Process boost periods are started and ended solely by z/OS operating system controlled events. The applicable events for Recovery Process boosts are limited to:

- HyperSwap™
- Coupling Facility data-sharing member recovery
- Coupling Facility structure recovery
- Sysplex partitioning

System Recovery Boost Stage 2 — WLM Support

- In contrast to shutdown boost, the boosted system will NOT be eliminated from the sysplex routing recommendation
- In contrast to startup boost, WLM's sysplex routing weight calculation will NOT be honor any additional capacity on the boosted system, and thus does not route additional work to the boosted system
- During RP boost, the algorithms of the sysplex routing services will handle the boosted system like not being boosted. This prevents additional work from other systems being attracted for a short period of time to the boosting system, while the recovery processing takes place – allowing the extra capacity to be applied meaningfully to recovery acceleration, not handling more workload

Further information

- **System Recovery Boost for the IBM z15 Whitepaper**
<http://ibm.biz/z15SRBWhitePaper>
- **System Recovery Boost**
http://publibz.boulder.ibm.com/zoslib/pdf/izsb100_v2r4.pdf
- **System z Mean Time to Recovery Best Practices**
<http://www.redbooks.ibm.com/redbooks/pdfs/sg247816.pdf>
- **Introducing IBM Z System Recovery Boost**
<http://www.redbooks.ibm.com/redpieces/pdfs/redp5563.pdf>

Other Enhancements and Important Information

OA558303 V2R4 / OA59890 V2R3 SRM ENHANCEMENT TO HANDLE DIAG 204 BUSY CONDITION

- In certain configurations, such as a low weighted LPAR or when an LPAR is being capped, it may take a considerable amount of time for PR/SM to complete a configuration change.
- This might delay SRM's request to PR/SM to obtain logical CPU utilization because PR/SM is busy.
- SRM is enhanced to handle PR/SM busy conditions when obtaining information on logical CPU utilization.

z/OS V2R5 SRM Contention Relief OA61240


- Timed algorithms are processed as an extension of any sysevent whenever they are due. In case the sysevent triggering the algorithm is running on a vertical low processor, the algorithm might not be executed and the SRM lock is held until the processor is dispatched again.
- SRM timed algorithms are only requested during a timer pop sysevent. This prevents SRM locks being held by vertical low processors.
- Added check if running under timer pop when requesting timed algorithms and skip the execution of algorithms that are due as an extension of a non-timer-DIE Sysevent
- No SRM locks are held due to algorithms requested that can not be executed. This leads to an overall better performance of the system.

XML Format WLM Service Definitions Recommendation

- Avoids particular problems in contrast to ISPF tables
- Since z/OS V2.3 XML-format is the default for "Save" and "Save as"
- Performance improvements in z/OS V2.4
- **Recommendation:**
Use the XML-format for your WLM service definition data sets

z/OS Workload Management

More Information



- z/OS WLM Homepage: <https://www.ibm.com/it-infrastructure/z/zos-workload-management>
- z/OS WLM GitHub: <https://github.com/IBM/IBM-Z-zOS/tree/master/zOS-WLM>
- z/OS MVS documentation
 - z/OS MVS Planning: Workload Management: [https://www-01.ibm.com/servers/resourcelink/svc00100.nsf/pages/zOSV2R3sc342662/\\$file/ieaw100_v2r3.pdf](https://www-01.ibm.com/servers/resourcelink/svc00100.nsf/pages/zOSV2R3sc342662/$file/ieaw100_v2r3.pdf)
 - z/OS MVS Programming: Workload Management Services: [https://www-01.ibm.com/servers/resourcelink/svc00100.nsf/pages/zOSV2R3sc342663/\\$file/ieaw200_v2r3.pdf](https://www-01.ibm.com/servers/resourcelink/svc00100.nsf/pages/zOSV2R3sc342663/$file/ieaw200_v2r3.pdf)
- IBM Redbooks publications:
 - System Programmer's Guide to: Workload Manager: <http://www.redbooks.ibm.com/abstracts/sg246472.html?Open>
 - Effective zSeries Performance Monitoring Using RMF: <http://www.redbooks.ibm.com/abstracts/sg246645.html?Open>
 - ABCs of z/OS System Programming Volume 12 <http://www.redbooks.ibm.com/abstracts/sg247621.html?Open>

Your feedback is important!

Submit a session evaluation for each session you attend:

www.share.org/evaluation

